Machine Learning Study Notes

Mind Map Style

Dr. Sarah Johnson • MIT - Computer Science Dept

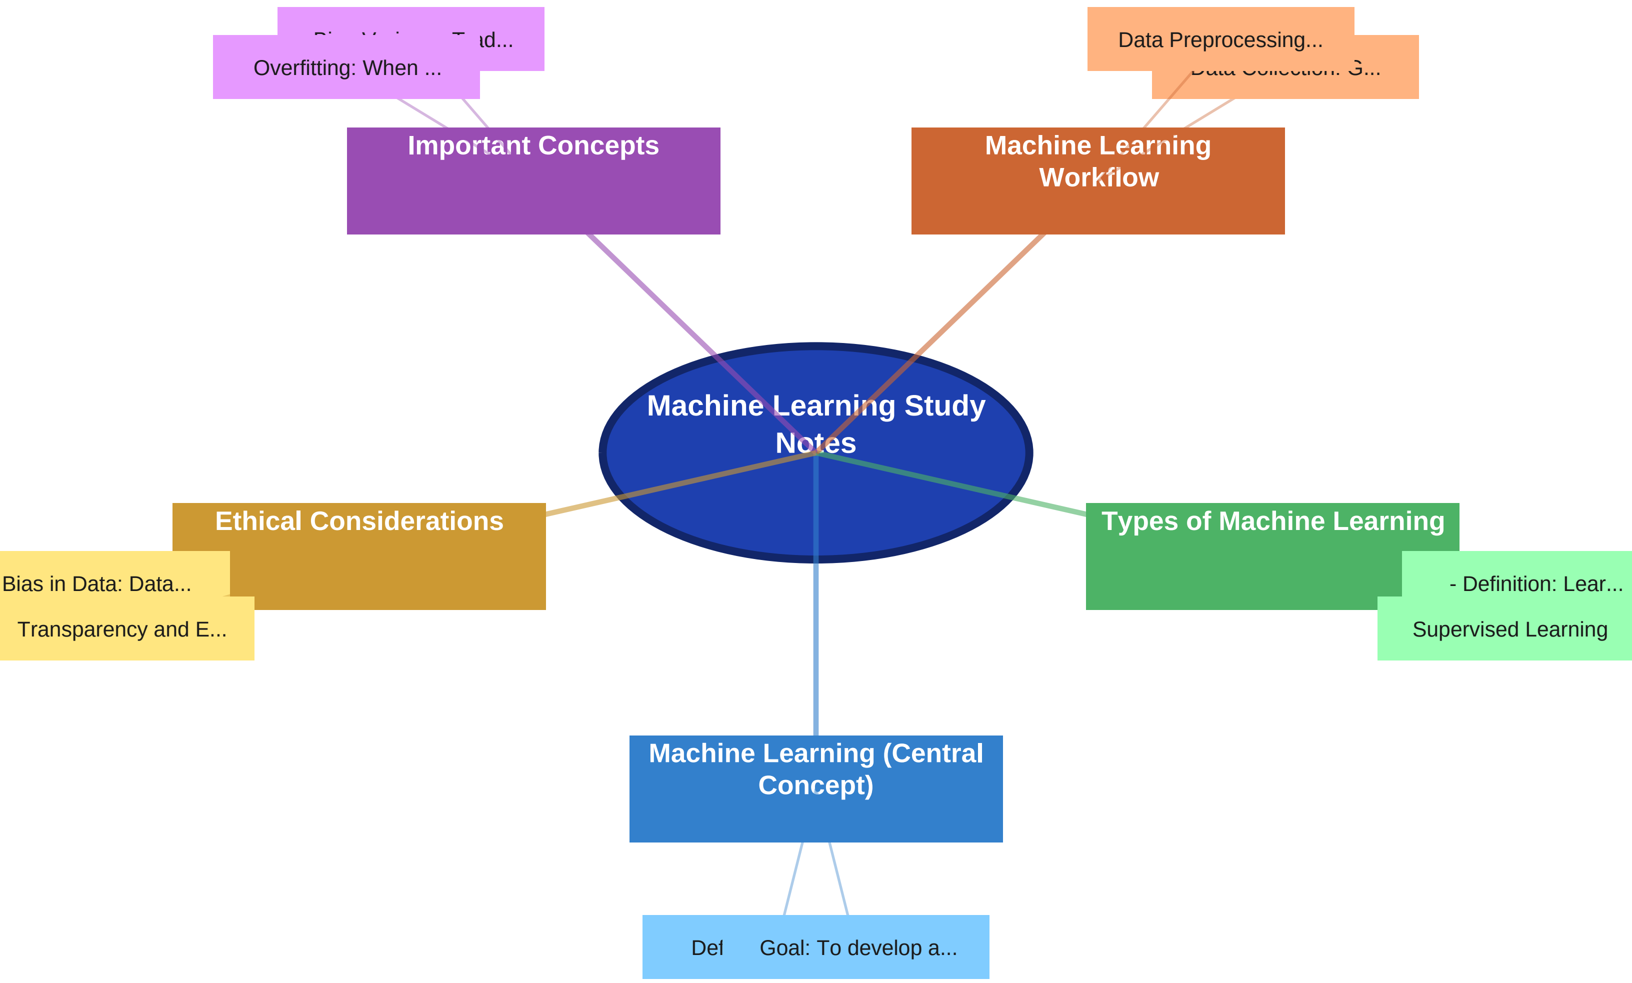

Detailed Notes

?? Notes

Machine Learning (Central Concept)

- Definition: Field of study that gives computers the ability to learn without being explicitly programmed.
- Goal: To develop algorithms that can learn from data and make predictions or decisions.

Types of Machine Learning

- Supervised Learning
- Definition: Learning from labeled data.
- Process: Algorithm learns a mapping function from input to output.
- Common Tasks:
- Classification: Predicting a categorical output (e.g., spam or not spam).
- Regression: Predicting a continuous output (e.g., house price).
- Examples:
- Linear Regression
- Logistic Regression
- Support Vector Machines (SVM)
- Decision Trees
- Random Forests
- Neural Networks
- Evaluation Metrics:
- Classification: Accuracy, Precision, Recall, F1-score, AUC-ROC.
- Regression: Mean Squared Error (MSE), Root Mean Squared Error (RMSE), R-squared.

- Unsupervised Learning
- Definition: Learning from unlabeled data.
- Process: Algorithm discovers patterns and structures in the data.
- Common Tasks:
- Clustering: Grouping similar data points together (e.g., customer segmentation).
- Dimensionality Reduction: Reducing the number of variables while preserving important information (e.g., Principal Component Analysis).
- Association Rule Learning: Discovering relationships between variables (e.g., market basket analysis).
- Examples:
- K-Means Clustering
- Hierarchical Clustering
- Principal Component Analysis (PCA)
- Association Rule Learning (Apriori Algorithm)
- Evaluation Metrics:
- Clustering: Silhouette score, Davies-Bouldin Index.

- Dimensionality Reduction: Explained Variance.

- Reinforcement Learning
- Definition: Learning to make a sequence of decisions in an environment to maximize a reward.
- Process: Agent interacts with an environment and learns through trial and error.
- Key Components:
- Agent: The learner.
- Environment: The world the agent interacts with.
- Action: What the agent can do.
- Reward: Feedback from the environment.
- State: The current situation of the agent.
- Examples:
- Q-Learning
- Deep Q-Networks (DQN)
- Policy Gradients
- Applications:
- Game playing (e.g., AlphaGo).
- Robotics.
- Recommendation systems.

- Semi-Supervised Learning
- Definition: A combination of supervised and unsupervised learning, using both labeled and unlabeled data.
- Motivation: Labeled data is often expensive or difficult to obtain, while unlabeled data is readily available.
- Common Techniques:
- Self-training
- Co-training

Machine Learning Workflow

- Data Collection: Gathering relevant data for the problem.
- Data Preprocessing: Cleaning, transforming, and preparing the data for the algorithm.
- Common Techniques:
- Handling missing values.
- Feature scaling (e.g., standardization, normalization).
- Encoding categorical variables (e.g., one-hot encoding).
- Outlier detection and removal.
- Feature Engineering: Creating new features or transforming existing features to improve model performance.
- Model Selection: Choosing the appropriate machine learning algorithm for the task.
- Model Training: Fitting the model to the training data.
- Model Evaluation: Assessing the performance of the model on the test data.
- Hyperparameter Tuning: Optimizing the model's hyperparameters to improve performance.
- Common Techniques:

- Grid Search
- Random Search
- Bayesian Optimization
- Model Deployment: Making the model available for use in a real-world application.
- Model Monitoring: Continuously monitoring the model's performance to ensure it remains accurate and reliable.

Important Concepts

- Bias-Variance Tradeoff: Balancing the model's ability to fit the training data (low bias) and generalize to unseen data (low variance).
- Overfitting: When a model performs well on the training data but poorly on the test data (high variance).
- Underfitting: When a model performs poorly on both the training and test data (high bias).
- Regularization: Techniques to prevent overfitting (e.g., L1 and L2 regularization).
- Cross-Validation: A technique for evaluating model performance by splitting the data into multiple folds and training and testing the model on different combinations of folds.
- Gradient Descent: An optimization algorithm used to find the minimum of a function (often used to train machine learning models).

Ethical Considerations

- Bias in Data: Data used to train models may contain biases, leading to unfair or discriminatory outcomes.
- Transparency and Explainability: Understanding how models make decisions is crucial for ensuring fairness and accountability.
- Privacy: Protecting the privacy of individuals whose data is used to train models.

?? Key Terms & Definitions

Machine Learning: A field of computer science that allows computer systems to learn from data without explicit programming.

Supervised Learning: A type of machine learning where the algorithm learns from labeled data.

Unsupervised Learning: A type of machine learning where the algorithm learns from unlabeled data.

Reinforcement Learning: A type of machine learning where an agent learns to make decisions in an environment to maximize a reward.

Classification: Predicting a categorical output.

Regression: Predicting a continuous output.

Clustering: Grouping similar data points together.

Dimensionality Reduction: Reducing the number of variables while preserving important information.

Feature Engineering: Creating new features or transforming existing features to improve model performance.

Overfitting: When a model performs well on the training data but poorly on the test data.

Underfitting: When a model performs poorly on both the training and test data.

Regularization: Techniques to prevent overfitting.

Cross-Validation: A technique for evaluating model performance by splitting the data into multiple folds.

Gradient Descent: An optimization algorithm used to find the minimum of a function.

?? Examples

Supervised Learning:

- Spam detection (classification): Classifying emails as spam or not spam based on their content.
- House price prediction (regression): Predicting the price of a house based on its features (e.g., size, location).

Unsupervised Learning:

- Customer segmentation (clustering): Grouping customers based on their purchasing behavior.
- Anomaly detection (clustering): Identifying unusual patterns in data, such as fraudulent transactions.

Reinforcement Learning:

- Game playing: Training an AI to play a game like chess or Go.
- Robotics: Training a robot to navigate a complex environment.

? Review Questions

1. What are the main types of machine learning and how do they differ?
2. Explain the bias-variance tradeoff.
3. What is overfitting and how can it be prevented?
4. Describe the machine learning workflow.
5. What are some ethical considerations in machine learning?

?? Summary

Machine learning is a field focused on enabling computers to learn from data without explicit programming. Key types include supervised learning (learning from labeled data), unsupervised learning (learning from unlabeled data), and reinforcement learning (learning through interaction with an environment). The machine learning workflow involves data collection, preprocessing, feature engineering, model selection, training, evaluation, and deployment. Important concepts include the bias-variance tradeoff, overfitting, and regularization. Ethical considerations regarding bias, transparency, and privacy are also crucial.